

Technical Notes on FieldWorks Send/Receive

July 27, 2017

Contents

1 Send/Receive Introduction.....	1
2 Getting started.....	4
2.1 Starting up Project (FLEX) Send/Receive.....	4
2.2 Starting up Lexicon (LIFT) Send/Receive in FLEX.	5
2.3 Starting up Send/Receive in WeSay.	6
2.4 Chorus Hub startup	6
2.4.1 FLEX Bridge 2.2 – Chorus Hub process	6
2.4.2 FLEX Bridge 2.3 – Chorus Hub service.....	7
2.5 LanguageDepot Internet startup.....	8
3 How it works, or why it doesn't.....	9
3.1 Lexicon examples	9
3.2 General concepts	10
3.3 Interlinear examples.....	11
3.4 WeSay/LIFT collaboration	11
3.5 FLEX/ParaText collaboration	12
3.6 Linked files	14
3.7 FieldWorks version.....	15
3.8 FieldWorks project name.....	15
4 Technical details	16
4.1 Chorus Hub and virtual machines.....	16
4.2 Using FieldWorks backups and Send/Receive	16
4.3 Restoring a FieldWorks backup.....	17
4.4 To switch a WeSay bridge user to another FLEX user.....	17
4.5 FLEX and WeSay compatibility issues	18

1 Send/Receive Introduction

FieldWorks Language Explorer (FLEX) provides two ways to collaborate with colleagues working on the same project. The first approach (starting in FieldWorks 7.3) includes the entire FieldWorks project: the full lexicon, grammar, interlinear texts, data notebook, scripture, lists, pictures, sound files, writing systems, and some configuration information. This is the approach that should be used whenever you are collaborating with other FLEX users. The second approach is via the Lexicon Interchange Format (LIFT). This approach should only be used if you are collaborating with WeSay users or some future program that can only use LIFT. LIFT only includes the lexicon, parts of grammar, writing systems, pictures, and sound files. The LIFT format does not cover as much detail in the lexicon as the FLEX project format, so the results are less satisfactory when using features of FLEX that WeSay does not use. There are also some extra issues with custom fields and lists.

With either approach, collaboration can be through any combination of the following

- Internet server (e.g., <http://languagedepot.org>) which requires some initial setup on the server that needs to be requested ahead of time.

- USB drive (no setup needed) that is passed around to users.
- Local Network using ChorusHub (requires setting up an instance of ChorusHub on one machine on a network).

If you are just collaborating between FLEx users, you would only use a FLEx repository (repo).

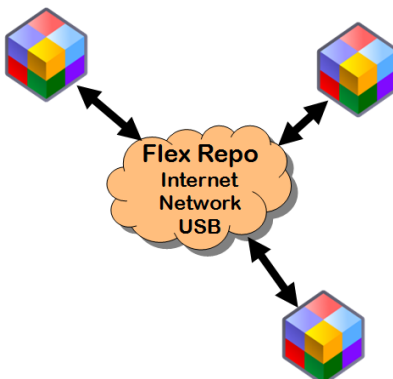


Fig. 1 Flex Collaboration

If you are just collaborating between WeSay users and one optional FLEx user, you would only use a LIFT repo.




Fig. 2 WeSay Collaboration

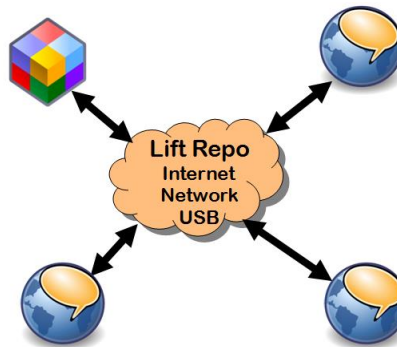



Fig 3. One Flex and WeSay Collaboration

If you are collaborating with multiple FLEx users and one or more WeSay users, then you will need a separate FLEx repo for collaboration between FLEx users, and a LIFT repo for collaborating between WeSay users and a FLEx user. It's critical (explained below) that only one FLEx user act as a 'bridge' between the two repos. This user would periodically sync with the LIFT repo and the FLEx repo. Other FLEx users would just sync with the FLEx repo, and WeSay users can only sync with the LIFT repo.



Multiple Flex and WeSay Collaboration

The FLEx Send/Receive menu has changed in different versions since FW7.3. This document was written for FLEx 8.0.5. Here is a brief summary of how the Send/Receive process works between two or more FLEx users.

One user starts the collaboration process from a master copy of the project. This user does an initial Send/Receive...Send this Project for the first time to store a copy of their project in a Mercurial repository (repo). Each colleague then gets a copy of the project from the repo using Send/Receive...Get Project from Colleague. After that, all colleagues periodically do Send/Receive...Send/Receive Project (with other FLEx users) (S/R) to keep their projects synchronized with their colleagues.

If colleagues are working in different parts of the project, the S/R will merge their changes without any conflicts or loss of data. If colleagues change the same piece of data (e.g., the definition of the same sense), during the S/R process the program will merge the changes the best it can (e.g., picking one of the two definitions), then add a conflict report warning the users about a change that was made that would be wise to review. If an undesirable change was made in a merge, the user will need to fix the appropriate data manually.

The current S/R process does not support different privileges for different users. Any user doing S/R can modify any part of the data. For an Internet server, a user can be given read-only permission to a repository, which allows them to get a copy of the project, but nothing they do will get back into the repo.

With each S/R operation, changes and new information are appended to the repo with a timestamp. Merges use a 3-way process that can generally tell whether the change was an addition by one user, a deletion by one user, or a modification by both users. Each user has a local (hidden) copy of the repo as it was the last time they did a S/R, so this helps the program to know what changed since their last S/R. FLEx does not provide a way to see what changed

during a S/R—it only shows merge conflicts that occurred. Some special programs outside of FLEEx let you see the history of changes in the repo, and it is possible to go back in time to an earlier version of the project, but this takes more advanced skills.

When using S/R with a project, you can still back up your project using FieldWorks backup, but you need to be very careful about using a FieldWorks restore, as this will likely cause all colleagues to lose their work since the backup was made.

2 Getting started

2.1 Starting up Project (FLEEx) Send/Receive

Before your first S/R you should make sure that you have a single project that can be considered your master project. If two users created their own FieldWorks project and added entries independently, the merge collaboration will not work. Also, if two users started with the same FieldWorks project, but have made independent changes, the merge collaboration will not work. This is discussed in more detail later. If you have multiple projects, you should find some way to merge these before doing your first S/R. Once you do the initial S/R of your project, all other collaborators need to delete or move (not just rename) any copies of the old project from their FieldWorks Projects directory and then start over by getting the master project from the new repo.

If you plan to use Chorus Hub or Internet options, you first need to set these up. See the following sections for details.

The first step is for one user to do an initial S/R from their master project. To do this, follow these steps.

1. Go to Send/Receive...Send this Project for the first time.
2. Optionally type a label in the Send/Receive Project dialog
3. a) Click USB Flash Drive or Chorus Hub if you are not using the Internet button. This will start the Send/Receive process to the specified device.
b) For Internet
 1. Click the Settings box in the lower right to bring up the Send/Receive Settings dialog.
 2. Make sure Server is set to LanguageDepot.org
 3. Fill in your LanguageDepot project id, user login and user password
 4. Click OK to close the Send/Receive Settings dialog.
 5. Now the Internet button should be enabled, so click that to start the Send/Receive process.

Every other user that wants to collaborate with this project then needs to do a one-time setup to get the project on their machine. To do this, follow these steps.

1. Make sure you have the same version of FieldWorks installed on your machine that was used for the master project.
2. Go to Send/Receive...Check for FLEEx Bridge Updates to see if there is a later version of FLEExBridge available. If there is, Click Install update and follow the default installation.
3. If you already have an older copy of this FLEEx project on your machine use File...Project Management...Delete Project if you can do this from another open project. Otherwise, delete the project folder in C:\ProgramData\SIL\FieldWorks\Projects. On XP this is in

c:\Documents and Settings\All Users\Application Data\SIL\FieldWorks\Projects. The default settings for Windows Explorer will not show this directory. You can still get to it by typing at least part of the path into the edit box at the top.

4. Start FLEEx. If the startup screen comes up, choose “Get project from a colleague”. Otherwise, go to Send/Receive...Get project from colleague. Either option will bring up a Receive project dialog.
5. Choose the location of the repo holding your project
 - a) Click USB Flash Drive or Chorus Hub, or
 - b) Click Internet. You’ll need to specify some additional information in the Get Project From Internet dialog before it will start.
 1. Make sure Server is set to LanguageDepot.org
 2. Fill in your LanguageDepot project id, user login and user password
 3. Type the project folder name you want to use on your computer (not the full path). The project folder will be created in your FieldWorks Projects directory.
 4. Click Download to start the download.

This process will create a new project folder in your FieldWorks Projects directory which will contain a copy of what was in the repo at that point.

Another way users can get started after the first user did the initial S/R is to copy the entire FieldWorks project directory from the first machine to the other machines. If you do this, the first time each user does a S/R, they should click the Settings link in the Send/Receive Project dialog and change the “Name to show in change history:” to the current user.

Once each user has a current copy of the project on their computer, they can periodically do Send/Receive...Project (with other FLEEx users) and click the desired location for synchronization. You can do another S/R to a different location if you want to keep several repos in sync.

2.2 Starting up Lexicon (LIFT) Send/Receive in FLEEx.

The master project for starting S/R can be either a FLEEx or a WeSay project. Note there are a few compatibility issues between FLEEx and WeSay. See Section 4.2 for more details.

If a FLEEx user is setting up the initial LIFT repo, use Send/Receive...Send this Lexicon for the first time (to WeSay) and then select the desired destination.

If a LIFT repo has already been set up by WeSay, and there is no FLEEx project for this language, then the FLEEx user needs to connect to the existing repo using Send/Receive...Get Project from Colleague. FLEEx will realize that the project is a LIFT repo, so it will create a new FLEEx project based on the LIFT repo.

It’s possible that you have been collaborating between FLEEx and WeSay in the past, but the connection may get broken. For example, if your FLEEx project directory gets deleted and you restore from a current FLEEx backup. (Note restoring a backup when using Send/Receive is usually dangerous. See Section 4.3 for more details.) To get reconnected to an existing LIFT repo, you need to use Send/Receive...Get Lexicon (WeSay) and Merge with this Project. This will do the initial sync to the LIFT repo without losing other information in your FLEEx project, such as interlinear texts. Before doing this command, inside your FLEEx language project folder, look for an OtherRepositories folder. If there is anything inside this folder, delete it first.

Once the initial connection has been established, then use Send/Receive...Lexicon (WeSay) whenever you want to sync with the LIFT repo.

2.3 Starting up Send/Receive in WeSay.

The master project for starting S/R can be either a FLEEx or a WeSay project. Note there are a few compatibility issues between FLEEx and WeSay. See Section 4.2 for more details.

If a WeSay user is setting up the initial LIFT repo, In the WeSay Home tab, click the Send/Receive button, and then choose the destination for the repo.

If a LIFT repo has already been set up by FLEEx or another WeSay user, and there is no WeSay project for this language on your machine, then you need to use the WeSay Configuration tool to get connected. In the WeSay Configuration Tool opening dialog, choose “Get from USB drive” or “Get from Internet”, choose the desired repo, then click Copy To Computer. This will create a WeSay project on your machine that will be in sync with the repo.

Once the initial connection has been established, then use Send/Receive in the WeSay Home tab whenever you want to sync with the LIFT repo.

2.4 Chorus Hub startup

Older versions of Chorus Hub used a process on the server computer. Recent versions use a service on the server computer. These two versions are described in the next two sections. It should be possible to use a mixture of versions of FLEEx Bridge and Chorus Hub without causing problems. But it is recommended that collaborators use current versions of software.

2.4.1 FLEEx Bridge 2.2 – Chorus Hub process

In this version of FLEEx Bridge, the Chorus Hub 2.4.266 program is installed in the FLEEx Bridge directory, and can be started on a single machine on the network. It is not a service, so it is stopped any time the user that started it logs off, or the machine goes into sleep mode. Only one instance of Chorus Hub can be running at one time on a network. If you try to start Chorus Hub when there is already an active Chorus Hub on the network, it will let you know that it can't be started because there is already a Chorus Hub running on the network. It uses a directory (c:\ChorusHub) on the machine that is running Chorus Hub to store repos for all users on the network. Anyone on the network has access to Chorus Hub whether it is running on your local machine or on another computer on the network. Chorus Hub supports repos from several different programs including FieldWorks, WeSay, and Bloom. Each program limits access to repos it can use.

From FW8.0.6 through FW8.1.4, chorushub.exe was installed in the FLEEx Bridge directory and could be started from there. C:\Program Files (x86)\SIL\FieldWorks 8\Installers\ChorusHubInstaller.msi was also copied to your machine during installation. This installer can be run on any machine, even without FieldWorks, to run the Chorus Hub 2.4.7 program on that machine. After installing the program, you can launch it from:
c:\Program Files (x86)\SIL\Chorus Hub\ChorusHub.exe (omit “(x86)” for 32-bit Windows)

When Chorus Hub is running, a Chorus Hub dialog is present that gives information about the program. Other users that start a Send/Receive operation on the network will have access to the Chorus Hub button in the Send/Receive dialog as long as some machine on the network is

running Chorus Hub. You'll have to wait a second or two for this button to be activated. Before you can close the Chorus Hub dialog, you must click the "Stop Chorus Hub" link at the bottom. The project directories will remain in the ChorusHub directory and can be used again by restarting Chorus Hub.

Initial startup of Chorus Hub may ask to permit hg.exe through the firewall. This should be granted. It may also need ChorusHub.exe clearance through the firewall. This seems to happen automatically with Windows Firewall, but with other firewalls you may need to allow these permissions before it works.

Chorus Hub 2.4 will give continuous warning messages, "this machine has more than one IP address", and may or may not work if it detects more than one IP address on your machine. This will happen if you have more than one network card, if you have wireless as well as wired connections, or if you are running virtual machines. In general, you should pick a machine to run Chorus Hub that has a single IP port. See Section 4.1 for further information on virtual machines.

2.4.2 FLEEx Bridge 2.3 – Chorus Hub service

Starting in FW8.2.0, FLEEx Bridge 2.3.5 or later is installed as part of the installation. This version uses a newer version of Mercurial (3.3.2) to solve some difficult bugs occasionally encountered with the older version. ChorusHub.exe is no longer installed in the FLEEx Bridge folder. A new ChorusHubInstaller.msi file is copied to c:\Program Files (x86)\SIL\FieldWorks 8\Installers. This is Chorus Hub 2.5.20 or later. This version installs a service called Chorus Hub Sharing Service on the computer instead of running a process. As a service, it will be available on the machine whenever it is running and is not dependent on a user being logged in.

Firewall: At least on some machines it is important to add c:\Program Files (x86)\SIL\Chorus Hub\mercurial\hg.exe to the inbound rules in Windows Firewall for Chorus Hub to work. If you get a S/R message, "abort: error: A connection attempt failed because the connected party did not properly respond after a period of time, or established connection failed because connected host has failed to respond" it probably indicates your firewall is blocking hg.exe.

When Chorus Hub is installed as a service, it can't easily warn you about another Chorus Hub already running on the network since services do not use dialogs. If there is more than one machine running Chorus Hub service on a network, as could happen in the past, FLEEx will pick one, seemingly the last one that was started. In the S/R dialogs Chorus Hub will identify the machine on which it is running. If for some reason you want to have Chorus Hub server installed on more than one machine, you should stop all services except the one you want to use. During installation, the service will be set to Automatic which means it starts whenever the machine is booted. To start or stop the service, you can type Services in the Start dialog, then click Services at the top of the menu. In this dialog, find Chorus Hub Sharing Service in the right pane. You can right-click this and choose to Start or Stop the service. With more than one installation, you should go to the Properties on this service and set the Startup type to Manual on all but your primary machine which should be set to Automatic so it will start up automatically when you reboot.

When you uninstall Chorus Hub through Programs and Features, it will automatically stop and uninstall the Chorus Hub service.

The Chorus Hub service will send messages to the Event Viewer when there are problems. To see these, type Event Viewer in the Start dialog, then click Event Viewer at the top of the menu

to open the event dialog. In the left of this dialog, open the Windows Logs node and click the Application node. Any messages from Chorus Hub will show up in the right pane.

If you install Chorus Hub service on a machine when another machine on the network is already running Chorus Hub, The installation takes place, but when it tries to start the service, it fails and gives a red error icon in the Application log. When you click this you'll see a message, "Only one ChorusHub can be run on a network but there is already one running on ..." The message will list the machine that is currently running Chorus Hub. The service will remain installed and set to Automatic. Thus when the computer is restarted it will try to start again and will succeed if the other Chorus Hub is no longer running. Otherwise it remains off.

Chorus Hub 2.5 service has the same limitations with multiple IP addresses on the Chorus Hub machine. It currently doesn't communicate this via the event log, but will cause a crash when someone tries to access Chorus Hub with an error report similar to this:

Msg: There was an error on the Chorus Hub Server, which was transmitted to the client.

.....

****Inner Exception:**

Msg: Could not connect to net.tcp://192.168.196.1:5912/. The connection attempt lasted for a time span of 00:00:20.9952009. TCP error code 10060: A connection attempt failed because the connected party did not properly respond after a period of time, or established connection failed because connected host has failed to respond 192.168.196.1:5912

See Section 4.1 for further information on how to solve this.

2.5 LanguageDepot Internet startup

LanguageDepot is an Internet server (<http://languagedepot.org>) that hosts repos for various programs including FieldWorks. As of FieldWorks 8.0.6, in order to use the Internet option, a repo needs to be set up on the server prior to your use.

The first step is to have user accounts set up for each user that will be working on your project. You can do this by going to the site and clicking the Register link at the top right. You'll need to fill in a Login (preferably firstName_lastName) and enter a Password. Do not use a password you use for other secure operations since the password is not encrypted as it is used. You will need to use this information when you set up Internet Send/Receive from your FieldWorks project. Fill out the other required fields and then click Submit.

Next, you need to e-mail admin@languagedepot.org to have the repo created. Let the administrator know whether you need a FLEx repo and/or a LIFT repo for your work, and give the vernacular language and ISO-639 code (see www.ethnologue.com) used for your project. Give the user Login name of the manager for this project. The administrator will then send back a project id that you'll need to use in Send/Receive setup.

A project can have one or more managers. A manager can sign in to <http://languagedepot.org> using their user login and password. They can then type in their project name in the search box and click on their project link to access information on their project repo. From the Settings tab and the nested Members tab, any manager can add other users as Manager, Contributor, or Observer. A manager or contributor can send and receive changes in the repo. An observer can get the project, but cannot make changes. If you click the outer Repository tab for your project, you can see information on each S/R to your repo. Clicking the number at the left will list files

that were modified for that S/R. Clicking a file link will attempt to show changes that were made to that file, although this may be very slow due to the size of many files and speed of the Internet. There are better ways to get detail on the repo if you need it. The Tortoise Hg Workbench will be described later.

3 How it works, or why it doesn't

For the most part, changes from all colleagues will be merged successfully. However, remember the program has definite limitations when merging results between users. If you don't consider these limitations, you may get results that are baffling and disappointing.

When users change different things in the project, the changes will merge without problem. However, when users modify the same thing, this results in a conflict that needs to be resolved. To keep the S/R process functioning quickly, we don't want to stop and ask the user about each conflict before finishing the S/R. So instead, the process picks one of the two conflicting changes and then adds a merge conflict report describing the conflict and what the program did to resolve it. Most common conflict descriptions will be clear to the user. However, there are some things that can change that are internal to the FieldWorks model, and there isn't a simple way to explain this to the user, so it may just show some underlying XML in the conflict report in these cases. Sometime after the S/R, the user can investigate the conflict reports and decide whether the program's guess was good or not, and fix any that were incorrect. At this point it requires a manual edit to fix any conflicts that were resolved incorrectly. The conflict report has a link that will try to take you to the spot that changed in FLE_x. Technically, if there is a conflict, the person doing the merge will win (except for deletions described later), but because it's not easy to determine which user will actually initiate the merge for a given object, it's best to assume that the winner is random. A conflict report is displayed after a S/R that results in conflicts. You can also see this report using Send/Receive...View Project Messages (for a FLE_x repo) or Send/Receive...View Lexicon Messages (for a LIFT repo).

Conflict information is stored in special XML files (*.ChorusNotes) that are merged along with the data. The user can mark conflicts as resolved so they don't normally show in the conflict dialog, but there isn't any way in the program to actually delete the conflicts. So if you get a huge number of conflicts, these conflict files can get very large.

3.1 Lexicon examples

Merging entries is not based on the headword, but on an internal unique identifier (id) that is assigned to an entry when it is created. The reason for this id is that headwords do not uniquely define entries since you can have homographs, and homographs are not unique because they can be renumbered or adjusted when switching to a different writing system, etc. Entries can also have the same headword but have a different morpheme type (prefix, suffix, root). Headwords can also be misspelled, but when corrected, we don't want the entry to be considered a different entry, especially when lexical relations and interlinear text depend on it. The computer needs something that is created once when an entry is created and it will never change for the life of the entry. The technical name for this is a Globally Unique Identifier (guid). Anything inside FLE_x that references this entry uses that id.

If I create an entry for 'house', and my colleague also creates an entry for 'house', each one will have a different id, so when we merge our projects using S/R, the result will have two entries for

'house'. Duplicate entries can be merged one at a time using the Tools...Merge with entry option, but it doesn't happen automatically. If two users take a list of words and both add them to the lexicon, after they are merged, as far as humans are concerned, there will be duplicates of every word. The computer knows they have different ids, so it thinks it did a great job of merging the entries. On the other hand, if one user adds words from A-M and another adds words from N-Z, and they merge their projects, everything will be great since the human and computer agree that these are different entries.

Likewise, the computer assigns a unique id for every sense that is created within an entry. Again, there isn't any other unique way a computer can identify a sense because some senses may have the same or missing gloss, definition, part of speech, etc. If I added a new sense for 'house' with a gloss 'political body', and my colleague did the same on his machine on the same entry, when we merge the projects, 'house' is going to have two identical senses from a human perspective, but the computer considers them different because they have different ids. These duplicate senses can be merged one at a time using the Merge Sense into menu option on a sense, but it is not automatic.

During the Send/Receive merge process, once the computer realizes two people changed the same entry or sense based on the id, then it will do the best it can at merging the changes within that 'object'. At least it knows it is working with the same entry or sense. If I add example sentences to senses, and my colleague adds semantic domains to senses, when we merge, the result will be exactly what we want. Likewise if I add Spanish glosses and my colleague adds French glosses, the merge will again be flawless.

If I change a gloss on a sense and my colleague doesn't do anything with that gloss, then when we merge, my change will be made in both projects. However, if I change a gloss for a sense, and my colleague changes the same gloss to something else, now when we merge we have a conflict, so the program will pick one change and reject the other, then display a merge conflict report after the merge.

3.2 General concepts

In FLEx, every 'object' has a unique id. Objects are defined in the underlying model of the data. Some examples of objects are lexical entries, senses, example sentences, etymology, pronunciations, allomorphs, grammatical info, lexical relations, reversal entries, interlinear texts, wordforms, wordform analyses, word glosses, parts of speech, notebook records, list items, etc. In WeSay only entries and senses have unique ids. The merge process basically adds any new objects, and merges the contents of existing objects with identical ids.

Now suppose I delete an entry or sense. As long as the other user does not modify this entry or sense, when we merge, the deletion I made will remain deleted. However, if my colleague changed something on that same entry or sense, the merge process doesn't want to potentially lose some important work, so the object I deleted will remain after the merge with the change my colleague made, and a merge conflict report will be added. So after the merge I may think the S/R messed up because the thing I deleted came back again.

When two or more people are working on a project, the longer the interval between S/R, the bigger the chance of having many merge conflicts. Naturally, if multiple users make a lot of changes to the same fields of the same entries or senses, a lot of merge conflicts will also result. So if you are working in the same areas, more frequent Send/Receives will be helpful. If you

plan to do extensive bulk editing or orthography changes, you can avoid a lot of conflicts if everyone does a S/R and then stops work until the extensive changes are made, then they all do another S/R before continuing their work to get the bulk changes.

3.3 Interlinear examples

Interlinear texts and charts are complex objects in FLEx with many sub-objects with ids and references pointing into the text via offsets. Even pasting text into the baseline will generate a whole new set of wordforms with different ids if the wordforms are new to the project. This is normally not a problem, but when using Send/Receive it presents major difficulties with merging if another user was working on the same text and will likely result in duplications of text and/or problems with interlinearization. So in order to avoid these potential merge problems, coordinate with your colleagues so that only one person works on a given text in any round of S/R operations, especially if one member makes any change to the baseline. Even adding or removing a segment break character (period, question mark, exclamation point, section sign) in the baseline will alter the underlying segment objects so that interlinearization of the affected segments from a colleague in the same S/R cycle may be lost.

Even when working in different texts and not modifying the baseline, two users doing interlinearization can create duplicate objects when working in the same S/R cycle. For example, if user A analyzes the word 'cat' for the first time, and user B also analyzes the word 'cat' in the same S/R cycle, when they merge, there will be two 'cat' entries in the lexicon as homographs, even though they have identical content. So one user will need to use the 'Merge with entry' option to combine these entries, then the 'Merge sense into' option to merge the two senses in the merged entry. Also, under Word Analyses, there will be two duplicate analyses for the wordform 'cat'. So one user will need to delete the duplicate analysis that isn't being used, or possibly use the 'Assign Analysis' option to move any analyses from one analysis to the other analysis.

Merges and deletions of analyses or word glosses in Word Analyses affect all interlinear texts that use the modified wordform. If another user makes use of the modified analysis or word gloss in the same S/R cycle, their analyses involving the modified wordform will likely be lost.

In summary, when working on interlinear texts, only one user should modify a given baseline in one S/R cycle. Any interlinearization that adds new entries, wordform analyses, or word glosses will produce duplicates if another user does similar interlinearization on the same wordform in the same S/R cycle. Any merges or deletions in the Word Analyses area will cause loss of analyses that were made to the modified wordforms in the same S/R cycle. Analyses that use existing entries, senses, and wordform analyses and glosses should survive S/R operations. Frequent use of S/R will minimize the chances of encountering these problems.

3.4 WeSay/LIFT collaboration

When collaborating between multiple FLEx users and one or more WeSay users, it's critical that only one FLEx user sync with the LIFT repo. This is because the logic for merging gets fuzzy when dealing with lexical relations, variants, complex forms, example sentences, etc. since these are not simple strings, but objects similar to entries or senses. In FLEx these each have unique ids, but the LIFT file used by WeSay does not record the unique id for these. Every time a S/R happens between a LIFT repo and FLEx, for any new objects other than entries and senses, FLEx

will give them a unique id. If multiple FLEx users sync with the same LIFT repo, each FLEx will have different ids for the same object, which would result in duplicate objects when the FLEx users merge via the FLEx repo.

Any program that uses LIFT files is supposed to read and modify what it understands, and leave the rest of the information intact for other programs that need it. FLEx uses a hidden LiftResidue field on several objects in the lexicon to store information that isn't in the FLEx model, so that it can be exported back to the LIFT file without being lost. Any time you import a LIFT file, every entry and sense will have a LiftResidue field added to hold this information. If you are only using FLEx to FLEx collaboration, the LiftResidue fields are excess clutter that can be deleted.

3.5 FLEx/ParaText collaboration

If a ParaText (PT) project is associated with a FLEx project, and you are collaborating with other users via PT Send/Receive (S/R), and the other users also have FLEx installed and are collaborating using FLEx S/R Project, and FLEx users are doing anything in the Texts & Words area, there are some cautions that need to be taken to avoid corrupting FLEx data.

A PT administrator is the only one that can associate a shared PT project with a FLEx project. Once associated, PT can access lexical data from FLEx, and FLEx can access current PT data in the FLEx Texts & Words area. The association is stored in the PT project settings file with the name of the associated FLEx project. These settings will go to colleagues during PT S/R. If the colleagues have a FLEx project with the same project name, they will automatically be connected with the PT project on their machine. If these users also use FLEx S/R Project to keep the FLEx projects in sync, there is potential for problems if they also use the Texts & Words area in FLEx.

Important! When you are using S/R with FLEx and PT, one user should load all of the data from PT, then go to Word Analyses. After doing this, he should S/R Project in FLEx. Then all other FLEx collaborators should do S/R Project to get all of this new data before they open the Choose Texts dialog in Texts & Words...Interlinear Texts. After this initial S/R cycle, it will work best if collaborators sync FLEx and PT projects at the same time. When new books are added to PT or significant changes are made, again, you should go through this same process with a single FLEx user initially getting the changes and then passing them around via S/R.

As with any interlinear text in FLEx where users are collaborating with S/R, you should avoid more than one user making changes in a given book during a single S/R cycle. For Scripture books, each section is treated as an interlinear text, so this caution would apply to each section in a scripture book.

Here is some technical detail to understand what is happening and what can go wrong.

When scripture is first loaded from PT by opening the Choose Texts dialog in FLEx and selecting scripture, for the books that are checked, it creates ScrBooks, ScrSections, ScrFootnotes, StText (for titles and headings), ScrTxtParas with ParseIsCurrent set to False, and Segments without Analyses. As part of the import, a checksum is made of the PT file and stored in ImportedChecksum of ScrBook. Once loaded, when a user opens the Choose Texts dialog, if the stored checksum matches the current PT checksum, then FLEx will not reimport that book. If the book is modified in PT, and a portion of that book is included in the current Text & Words area, when the user goes to the Choose Texts dialog, FLEx will recognize that the PT book has

changed, so it reimports the book, trying to maintain any interlinearization that has been done on that book. These checksums are included in FLEEx S/R.

When you click a scripture text in the Texts pane (actually a ScrSection, ScrFootnote, or StText for titles and headings) that text is parsed creating WfiWordforms and PunctuationForms as needed, setting default Analyses on Segments, and setting ParseIsCurrent on ScrTxtParas to True.

If you go to the Word Analyses (or Concordance) tool, it will parse all texts and scripture that is currently imported, creating WfiWordforms and PunctuationForms as needed, setting default Analyses on all Segments used in those texts, and setting ParseIsCurrent on all ScrTxtParas to True.

All of the objects above, except Analyses which are initially pointers to WfiWordforms or PunctuationForms, have unique IDs. So if two FLEEx users import the same books in the same S/R cycle, there will be a lot of duplications that would be hard to clean up after a S/R merge. FLEEx does not provide a way to delete duplicate scripture books or sections that get imported in this way. Although wordforms have unique IDs, they will actually merge during S/R if the form is the same, so that helps some.

The duplicated portions can result in the load process failing. This will continue until the book is cleaned up enough for the import process to work. With Translation Editor (TE) you could run a verse chapter check on a bad book which will usually show serious problems in the data. If you clean these up enough in TE, you should get to a point where the import will work again.

If something is messed up too badly, it may be necessary to delete a book entirely. This is best done using FDOBrower. This is a low-level utility program that is in c:\Program Files (x86)\SIL\FieldWorks 8 directory. You should not have FLEEx or TE open when using this program. To delete a book of scripture, you can type FDOBrower at the Windows start command and it should find the program.

Once FDOBrower is open, go to File...Open Language Project. In the open dialog, you'll have to navigate to your project *.fwdata file. Project directories are normally under c:\ProgramData\SIL\FieldWorks\Projects\. After the project is open, you'll see two other tabs. Click the LangProj tab then scroll down to TranslatedScriptureOA and click the + to the left to expand the node. Under the TranslatedScriptureOA node, open the ScriptureBooksOS node. You'll see a node for each book that has been imported from PT. To delete a book, right-click the book and choose Delete. When done, do File...Save Current Language Project and then close FDOBrower. FDOBrower does autosaves the same as FLEEx.

Various things can go wrong when users sync FLEEx and PT at different times and the user goes to the Choose Texts dialog. For example, if user A imports scripture and then does FLEEx S/R, his import and checksum will go to user B. But if PT is not in sync with users A and B and user B does FLEEx S/R and opens the Choose Texts dialog, they will reimport the scripture from B's version of PT data, and conflicts will likely happen the next time they do FLEEx S/R with user A. Currently, the FLEEx merge conflicts for scripture typically give information that is difficult even for technical people to understand. Also, the merge conflict hotlinks into FLEEx do not work for scripture at this point. So it's best if colleagues work in a way that reduces or eliminates scripture merge conflicts.

By following the Important paragraph above, most conflicts will be avoided. If one user imports books from PT and then does FLEEx S/R and PT S/R, and another colleague does FLEEx S/R and PT S/R, their FLEEx and PT projects will normally be in sync so that nothing will be reloaded by going to the Choose Texts dialog. However, if multiple users are importing scripture in one S/R cycle, the FLEEx S/R merging capability may have trouble sorting everything out.

Interlinearization on Scripture is stored with each paragraph in a section. The entire section is loaded into FLEs as an interlinear text. When FLEEx imports from PT, it will update the baseline text and then try to maintain the interlinearization of the revised text. If you make a lot of changes to paragraphs, and add and remove sections heads in PT, there is a good possibility that FLEEx will not be able to keep up with the changes which means interlinearization may be lost. Also, if you delete a scripture book as described above to get past a messed up book, you'll lose all interlinearization in that book.

When we talk about interlinearization being lost, it doesn't mean everything is lost. There are 3 places that are affected by interlinearization. First, you may create a wordform in the wordform inventory and add analyses to the wordform. Second, these analyses are connected to entries and senses in the lexicon, which may involve creating entries or senses in the lexicon. Third, in the interlinear text itself, for each wordform in a paragraph you are linking to one of the analyses in the wordform inventory. When we talk about losing interlinearization, what we are actually losing is the pointers to specific analysis in wordforms which means it will default to the wordform. So the actual work in creating wordforms, analyses, entries, and senses are not lost and are still available when doing future interlinearization. So although you do lose work by deleting a scripture book (or interlinear text), recovering should be much faster because everything you previously created in the lexicon and wordform inventory is still there and you simply have to link the wordforms again. FLEEx will propose analyses for the entire text, so the proposals can either be approved immediately, or adjusted to get the correct analysis.

The process of associating a PT project with a FLEEx project was designed for one FLEEx project being associated with one PT project. If you happen to have multiple PT projects for a single FLEEx project, you should stick to associating a single PT project to the FLEEx project. If you switch the associations around between different PT projects, it will likely confuse the process resulting in corrupted Scripture and loss of interlinearization.

3.6 Linked files

Pictures, sound files, and other linked files are included in the S/R process as long as they are stored under the default LinkedFiles directory inside the FieldWorks project directory. FLEEx allows you to use an external directory for linked files which is advantageous if you have multiple projects that refer to a master set of pictures and sound files. But if you have chosen this approach, they will not be included in S/R. Also, because repo size and time for S/R, especially to the Internet, can become too great if you use high resolution pictures, sound files, movies, etc., the S/R process currently limits files to 1 Mb, and only accepts certain file extensions. For images, it accepts these extensions: bmp, jpg, jpeg, gif, png, tif, tiff, ico, wmf, pcx, and cgm. For audio, it accepts these extensions: wav, snd, au, aif, aifc, aiff, wma, and mp3. Anything that doesn't meet these requirements is skipped during S/R. A warning message will be given in the S/R log if files are greater than 1 Mb. Files with doc and txt extensions are included, but mp4 files are not.

3.7 FieldWorks version

Each version of FLEEx works with a single underlying data model. Any version of FLEEx can share projects from other versions of FLEEx as long as the data model is the same. When a user upgrades FLEEx to a version that uses a different data model, in any project he opens, data will be migrated to the new model. At this point users with older versions of FLEEx will no longer be able to open the project until they upgrade to a FLEEx version that supports that data model.

When using S/R for a FLEEx project, all collaborators should use the same version of FLEEx, or at least versions that use the same data model. You can see the model version in a FieldWorks project directory involved in S/R by opening the file, `FLEExProject.ModelVersion` in a text editor. It contains a version, such as 7000070). For any project, the data model can also be seen in the second line of the `fwdata` XML file. Recent data models and associated versions of FLEEx are

7000068 FW8.0.4 – FW8.2.9

7000069 FW8.3.1

7000070 FW8.3.2 –

FLEEx stores data in the mercurial repo in a unique branch for each data model. A repo can have different branches, and support users in each branch, but data will not be merged or made available between branches until the last colleague upgrades to the most recent version.

Thus if users A, B, C, and D are collaborating in version FW8.3.1 everything works fine. If user A upgrades to FW8.3.2 and does a S/R the repo now has two branches. If user B also upgrades to FW8.3.2 and does S/R, A and B can now collaborate with each other, and C and D can collaborate with each other, but they can't see any work done by the other two users. If C upgrades, his work as well as the work merged with user D will merge with A and B so that all three of them can collaborate. Once D upgrades and does S/R, the two branches will be merged so that everyone will again be working on a single branch.

For this reason it is best for collaborators to upgrade about the same time so that they can continue to see each other's work.

3.8 FieldWorks project name

Repos also have a unique id that is assigned when it is first used. The id stays with the repo when you do a S/R to different locations, or copy the repo. This is used to ensure that you are really working with the correct repo. During a S/R the program uses this repo id to find the project in the FieldWorks Projects directory. If one user changes the name of their FieldWorks project, they can still S/R with the same repo because the repo id hasn't changed. For this reason, if you want to receive a new version of the repo, you can't just rename the old folder or give it a different name in the Receive dialog. In order to do this, you need to actually move the old project folder outside the Projects folder, or move it inside a subfolder inside the Projects folder so that it won't cause a conflict. This also means that different colleagues can specify their own name to a project when it is received the first time, although this could be rather confusing.

4 Technical details

4.1 Chorus Hub and virtual machines

Chorus Hub 2.4 will give continuous warning messages, “this machine has more than one IP address”, and may or may not work if it detects more than one IP address on your machine. Chorus Hub 2.5 is unable to display these warning messages, but has the same limitations. Instead, when a user attempts to access Chorus Hub 2.5 under these conditions, they will get a crash with a message similar to the following:

Msg: There was an error on the Chorus Hub Server, which was transmitted to the client.

.....

****Inner Exception:**

Msg: Could not connect to net.tcp://192.168.196.1:5912/. The connection attempt lasted for a time span of 00:00:20.9952009. TCP error code 10060: A connection attempt failed because the connected party did not properly respond after a period of time, or established connection failed because connected host has failed to respond 192.168.196.1:5912

If the problem is due to a virtual machine, and you really need to use Chorus Hub on the virtual machine or on the machine that is running a virtual machine, there are ways to make this possible. Here are some notes for using Oracle VM VirtualBox on Windows 7. There are probably similar things that can be done with other virtual machines.

Virtual Box defaults to creating an IP address for the virtual machine. The default attachment is NAT which is a separate IP address that is not available outside the virtual machine. If you need to run Chorus Hub in the virtual machine, you can change this setting to Bridged Adapter. In this mode it uses the same IP as your main machine. So when you run Chorus Hub inside the virtual machine, users outside the virtual machine will be able to use it.

If you want to run Chorus Hub on a machine that is using a virtual machine, you’ll need to disable the IP address for the virtual machine. You can do this using Control Panel, Networking and Sharing. Click “Change adapter settings”, then right-click the Virtual Box Host-Only Network and choose disable. Now Chorus Hub will run on the main machine without complaining about extra IP addresses. You can still access the Chorus Hub running on your main machine from within the virtual machine.

4.2 Using FieldWorks backups and Send/Receive

Making FieldWorks backups is not as critical when using Send/Receive Project because the S/R process is storing a type of backup on the repo. If your machine should fail or something happens that makes your version of FLEEx unusable, you can delete your FLEEx project directory and start over again by using Get Project from colleague. This will restore your project to the current state of the repo.

If you are only using S/R Lexicon, then the repo only stores the lexical data, but none of the other data that may be in your FLEEx project. So in this case, you should definitely do regular FieldWorks backups to prevent data loss in case of some emergency.

While making FieldWorks backups is perfectly safe when using S/R, you normally should not do a restore from a backup when using S/R. For one thing, if you restore your project from a backup

and then do S/R, it will basically set all other users back to your restored state, thus cancelling all work they have done since the state of the project when the backup was made.

Also, a FieldWorks backup does not include the LIFT or FLEEx repo that is normally inside your project directory. So if your project directory gets deleted, when you restore from a backup, you are now disconnected with both the LIFT and FLEEx repos.

If you are using S/R Project, instead of restoring from a FieldWorks backup, you should delete your directory and use Get Project from colleague to restore your project. Then it would be in sync with the repo. If you have made a FLEEx backup since your last S/R, the next section describes how you can restore from this.

If you are using S/R Lexicon, and you need to restore your project from a FieldWorks backup, then you should use Get Lexicon (WeSay) and Merge with this Project. But keep in mind that this will probably cause a loss of the WeSay work since the time the FieldWorks backup was made.

4.3 Restoring a FieldWorks backup

Restoring from a backup is usually not a good thing when using S/R, so to prevent accidentally losing a lot of data, the program currently refuses to let you restore to the project, and forces you to restore to a different name. Say your original project is QQQ and you restore to QQQ-01.

There isn't a direct way to get QQQ-01 hooked back up to S/R. The only way you can get back to S/R with the restored data is to copy the QQQ-01.fwdata file from the QQQ-01 project to the QQQ project and rename it QQQ.fwdata, replacing the bad QQQ.fwdata and then continue S/R from the QQQ project. This will be safe to do IF the backup from which you restored was made since your last S/R. Otherwise this would likely cause loss of data for everyone on the project.

The normal way to recover if you mess up the project in a way that you don't want to pass on is to delete your QQQ project and then do Get Project From Colleague to get back to a normal S/R project. Note, if you have QQQ-01 in your Projects directory, you will not be able to do Get Project From Colleague because it searches all projects to see if any have the same unique project id stored in the fwdata file. This project id will be the same if you restore a project from a backup or rename the project. So you need to make sure there is no project in the projects directory that was made at some point from the project you want to download.

If you are going to do something potentially dangerous to a project that is connected with S/R, one approach is to first make a copy of the fwdata file in the project directory using Windows Explorer. If you need to restore, just delete the fwdata file and rename the copy back to the original name. But again, this should only be done if you haven't done S/R since the time you made the backup.

4.4 To switch a WeSay bridge user to another FLEEx user

Remember that only one FLEEx user should provide a S/R bridge between FLEEx and WeSay users. Assume FLEEx users A and B are using S/R Project to keep in sync, while user A is also doing S/R Lexicon to sync with WeSay users. Now they would like to switch the WeSay bridge operation to user B instead of A. If this process is not done correctly, you can duplicate senses, examples, etc. To avoid these problems, use the following process.

1. All WeSay users do S/R and stop work.

2. All Flex users do S/R Project and stop work.
3. User A does S/R Lexicon to get the latest changes from WeSay users.
4. User A does S/R Project to send these merged changes to the FLEEx repo.
5. User B does S/R Project to get the latest from the FLEEx repo.
6. User B does Send/Receive...Get Lexicon and merge with this project to get started as the bridge.
7. User B does S/R Project to push any changes to the FLEEx repo.
8. At this point all users can do S/R and then continue working.

At this point user B would be the only one that does S/R Lexicon to function as the bridge between WeSay and FLEEx users.

When done, user A should delete the *_LIFT directory inside their FieldWorks project folder under the OtherRepositories folder to make sure they don't accidentally do S/R Lexicon.

The reason for needing these steps is that S/R Get Lexicon and merge this project uses the 3rd LIFT import option that keeps objects from both users when there is a conflict rather than one user overwriting the other. The above process makes sure there will be no merge conflicts during the process of switching the bridge.

4.5 FLEEx and WeSay compatibility issues

There are some compatibility problems with writing systems between WeSay and FLEEx for non-standard ISO codes. Lists are maintained in different ways between the two systems and are not included in S/R. Custom fields are also stored in different ways.

WeSay is heavily biased toward lexeme forms and definitions, and this bias cannot be removed in configuration. In WeSay 1.5.38 you can check Gloss in the dictionary configuration, but it doesn't show the field in the editing mode unless it currently has data. So when adding a new entry or editing an existing entry without gloss, you have to first click "Show Uncommon Fields" in the lower right to open the Gloss field. When adding a new sense in WeSay, you must enter a Meaning (definition) before you can add a gloss. But in the Semantic Domains tab, it does not give any option other than Meaning, which is definition.

If I create an entry in FLEEx with only a gloss, and Glosses are enabled in WeSay Configuration, it shows my gloss in both the Meaning (definition) and Gloss fields. If I edit the gloss in WeSay, only the gloss is changed. Likewise if I edit the definition in WeSay, only the definition is changed. If I add an example in WeSay, the definition will now contain the same content as the gloss even though I didn't touch either one. If you don't edit anything in the entry, WeSay doesn't add the definition to the data even though it shows the gloss in the Meaning field.

If Gloss is not checked in WeSay Configuration, if an entry comes from FLEEx with only a gloss, it shows the gloss content in the Meaning field in WeSay and if you edit it, the edited field becomes the Definition and the Gloss is unchanged. If you don't edit the meaning, but make some other change in the entry (e.g., add an example) the definition becomes real and is a copy of the gloss.

WeSay stores lexeme forms in the Word field. It is possible to configure WeSay to show Citation Forms, but this is limited. It is possible to click "Show Uncommon Fields" to add a Citation Form without filling in the Word (lexeme form) field. But in the list of entries, it shows a blank instead of showing the Citation Form. And in the Semantic Domains tab you can only add a Word (lexeme form).

If you plan to use FLEx and WeSay with Send/Receive during a rapid word collection process, you need to use lexeme form and definition in both programs. If you want to do anything else, then you should only use FLEx during the process and configure it for a simple view that only shows what you want.